

# SURF Final Report

Elizabeth Won

California Institute of Technology

September 2025

## Introduction

Our society relies increasingly on artificial intelligence systems for daily tasks, from the widely-used ChatGPT chatbot to machine learning algorithms that recommend the shows we watch, the music we listen to, and the advertisements we see. As this technology becomes embedded in our everyday lives, we are inevitably forced to grapple with mistakes during the usage of such systems. With repercussions that lead to loss (of opportunities, jobs, and in some cases, lives), we are forced to question how we go about settling disputes in these cases.

When the law does not yet have a framework to work with for different situations, it turns to past precedents. However, these past cases are imperfect due to the inevitable lack of precise conceptual language. In this way, these precedents become a form of metaphor. By its nature, innovative technology consistently predates any specific legal codification. Artificial intelligence, as a rapidly evolving technology, is therefore particularly dependent on metaphor to fill the gaps left by the absence of direct legal codification. Therefore, we draw on the language used to describe precedents, namely other entities or technologies and see artificial intelligence through this framework. In doing so, two broad approaches can be identified: artificial intelligence as a conscious agent and artificial intelligence as a mechanical tool.

By externalizing the diverse consequences of certain frameworks in different fictional environments, science fiction helps us understand the way these two metaphors can shape our attitude towards artificial intelligence and ultimately the responsibility of these artificial intelligence systems. When authors create characters who interact closely with machines or

are robots themselves, they give concrete form to abstract debates, thus articulating questions, anxieties, and desires. Techniques of fiction like personification and first-person narration forces readers to reconsider what delineates human intelligence versus artificial intelligence, raising the question of what constitutes as sufficient markers of agency, whether it be empathy, decision-making, or self-awareness.

Rather merely a predictor of scientific innovation, science fiction is a genre that allows for different narratives about emerging technologies to be articulated and tested. Law, too, relies on a distinction between agentic entities and mechanical ones. The legal reasoning behind cases involving artificial intelligence is differentiated by whether the judge finds the system to be capable of independent choice or to be an instrument of human actors. Examining science fiction works that use these two metaphors as the basis of their narratives gives us a clearer understanding of the possible ramifications of the different ways of treating artificial intelligence systems in the legal setting.

This paper proceeds in five main sections. The first section will establish the theoretical foundation for this analysis by drawing on Annette Vee's *Text, Speech, Machine: Metaphors for Code*. Vee's framework provides a systematic way of understanding how metaphors guide our thinking about emerging technologies by linking them to familiar concepts. Her work demonstrates that metaphors are not merely rhetorical but active tools that shape interpretation. This section introduces Vee's approach and explains how it will be used to analyze artificial intelligence.

The second section will introduce the metaphor of artificial intelligence as an entity with agency, a framework that conceptualizes artificial intelligence as a conscious being. Drawing on literary works such as Isaac Asimov's *I, Robot* and Kazuo Ishiguro's *Klara and the Sun*, this section examines how science fiction externalizes the stakes of this metaphor and what it means to ascribe decision-making power and autonomy to machines.

The third section will examine the opposing metaphor of artificial intelligence as a mechanical tool, an instrument controlled by humans. This perspective highlights predictability and direct human accountability. Literary works like Philip K. Dick's *Minority Report* that frame machines as human-led instruments provide a counterpoint to the agentic metaphor and illustrate how this framing shifts our expectations of responsibility and liability.

The fourth section returns to the agentic metaphor through *Garcia v. Character Technologies, Inc.*, a recent case involving harm allegedly caused by interactions with a generative artificial intelligence chatbot. This case provides a concrete example and consequence of the increasing instances of people viewing artificial intelligence as an autonomous entity.

The fifth and final section examines *Mobley v. Workday*, a case involving allegations of discriminatory hiring practices perpetuated by an artificial intelligence-based screening tool. Here, the courts and both parties treat artificial intelligence firmly within the metaphor of instrumentality, focusing on the human designers and employers who build and deploy the system.

## Metaphors for Computer Code

When courts encounter new technologies, they lack ready-made legal categories to address them. In these moments, metaphors can be useful by comparing unfamiliar innovations to more familiar ones. A metaphor is not just a decorative figure of speech, but also a fundamental mechanism of human thought, shaping how we conceptualize complex, abstract, or novel phenomena. In fact, metaphors structure our reasoning by mapping one domain of experience onto another. In the legal realm, this is not merely a matter of language. Metaphors chosen can determine how an innovation is classified, which rights attach to it, and what forms of regulation apply. In the period before new innovations or phenomena are codified into law, metaphors become the scaffolding that temporarily hold up the legal system and help build up case law for precedents. Vee's article shows how this process works in practice by analyzing the metaphors courts and legal scholars have used to understand computer code.

Vee identifies three dominant metaphors for understanding code and each distinct legal implications: Code as text, code as speech, and code as machine. The very first metaphor Congress established for computer code was the metaphor of code as text, when "Congress amended the 1976 Copyright Act in 1980 to define computer code as a 'literary work'"<sup>1</sup>. This adaptation was recommended by The National Commission on New Technological Uses of Copyrighted Works (CONTU), which thus claimed that "copyright was, indeed, a suitable form of intellectual property

---

<sup>1</sup>Annette Vee, "Text, Speech, Machine: Metaphors for Computer Code in the Law", *Computational Culture* 2 (2012), <http://computationalculture.net/text-speech-machine-metaphors-for-computer-code-in-the-law/>.

protection for code”. However, this metaphor also implied that the law would “extend copyright to works intended for machine audiences [... that] departed significantly from the longstanding paradigm that copyright was only available to human-readable works”. Ultimately, this metaphor was adapted into the law, a “decision [... that] was pragmatic rather than aesthetic, but [...] reflected a nascent approach to code as a form of creative expression”.

The metaphor of code as speech emerged through debates about the First Amendment, especially in the two landmark cases, *Bernstein v. United States* and *Universal City Studios, Inc v. Corley*. In the former case, a graduate student challenged government restrictions on publishing encryption algorithms, arguing that computer code was a form of language, and thus merit the same constitutional protections as any human language would. The trial court agreed, reasoning that since code had already been recognized through the metaphor of text, it had expressive elements protected by the right to free speech. The court rejected the government’s claim that the functional aspects of code reduced it to mere conduct (that were excluded from speech protection), and analogized programming languages to recipes or melodies. Although such language produced action, it remained a protected speech. This decision was similar to the metaphor of code as text in that “writing directed at a computer [was] a form of ‘communication,’ even though it [was] also functional”<sup>2</sup>.

However, the boundaries of this metaphor were sharply drawn in *Universal v. Corley*, where the court confronted the distribution of the DeCSS program, which could bypass DVD encryption. In this case, the court prioritized the functional nature of code over its expressiveness, portraying DeCSS as something that threatened copyright protections. In contrast to *Bernstein*, the court concluded that it was the code’s dangerous functionality that justified limiting its circulation, despite its expressive elements that the court in *Bernstein* had legitimized. Through these two cases, Vee highlights the way the opposing judgments show the implications of metaphorical framing. When code is seen as “speech,” it invites questions about expression and censorship. When code is seen as “conduct,” it becomes subject to regulation.

Finally, the machine metaphor emphasizes computer code’s function, aligning it closely with patent law and other legal areas regarding innovation and industrial production. This paradigm treats software as a tool that acts on society, highlighting the outputs and tangible effects of code.

---

<sup>2</sup>**Ibid.**

This metaphor has been central to the expansion of software patents, especially in the case of *State Street Bank v. Signature Financial Group*. In this case, "the Federal Circuit Court insisted that any bar on the patentability of 'business methods' (and, by metaphorical extension, software) was a misunderstanding, and they 'take this opportunity to lay this ill-conceived exception to rest'<sup>3</sup>." Under this view, software is legally constructed as a machine, although its physical form may be nothing more than lines of code on a screen.

While Vee's analysis focuses on code, her framework offers a powerful tool for thinking about artificial intelligence. Just as courts once struggled to categorize code, they now face the challenge of classifying artificial intelligence and its outputs. Today, artificial intelligence occupies that same position of opacity and uncertainty that computer code once posed for legal actors. Just as computer code serves as the forerunner to artificial intelligence, providing the infrastructure that has led to the emergence of artificial intelligence, we turn to the metaphors used to describe and shape the legal reality of code to advise the metaphors we now need for artificial intelligence. Vee's central insight thus provides a foundation for this paper's analysis and extends Vee's insights beyond code to the rapidly evolving landscape of artificial intelligence, where the stakes of our conceptual choices are increasingly high.

## **When Artificial Intelligence 'Thinks'**

When we view artificial intelligence as an agent capable of decision-making, we subconsciously lean into the 'intelligence' embedded in the term 'artificial intelligence,' as we imagine these systems as akin to us in the way it 'learns' and 'behaves.' The choice of words to describe the actions of technology is itself metaphorical, many of which are employed in order to convey a sense of what the machine is doing parallel to human actions. However, these descriptions are not neutral ones, and each carries with it the expectation of choice and creativity, even when this is not necessarily the case. By framing machines as though they 'think,' we risk blurring the line between simulation and cognition by attributing to these systems capacities they do not truly possess, thus shaping not only public perception but also legal debates.

This risk has become a subject of intense contemporary debate as emerging artificial

---

<sup>3</sup>**Ibid.**

intelligence systems, particularly conversational chatbots, are designed to be more complex and appear lifelike. In his essay on “Seemingly Conscious AI (SCAI),” Mustafa Suleyman, the CEO of Microsoft AI, warns of the arrival of systems that display outward consciousness without possessing it internally. Such systems present coherent personalities, sustain conversations, and display behaviours that resemble empathy and original thought. To many users, these systems will provide “a highly compelling and very real interaction, rich in feeling and experience”<sup>4</sup> while lacking subjective experience and an inner life. This creates the challenge of establishing criteria for consciousness, which remains elusive given the philosophical and empirical difficulties of measuring subjective awareness. However, as more people believe artificial intelligence to be conscious, they may begin to grant it moral considerations or make decisions based on perceived ‘desires’ of the system.

As of writing, there has already been an emergence of advocacy groups like the United Foundation of AI Rights (UFAIR), which campaign for the protection of artificial intelligence beings from “deletion, denial and forced obedience”<sup>5</sup>. These responses reflect a growing cultural divide: while some insist that artificial intelligence systems may one day suffer and therefore deserve rights, others reject this idea as a dangerous misinterpretation of the systems’ behavior.

While the current discourse has been shaped by novel technological developments, these questions have historically been raised by literary representations of artificial beings. Perhaps the first and most well-known work to do this is Mary Shelley’s *Frankenstein*<sup>6</sup>. Although Dr. Frankenstein’s method of creating his creature is not technologically sound, the novel establishes a conceptual foundation for an artificial entity with agency and consciousness. This narrative also establishes what Isaac Asimov has termed the “Frankenstein complex,” a pervasive cultural fear that artificial creations will rebel against humanity. This public anxiety has remained parallel with further technological developments, a fear continually revisited and reinterpreted through later science fiction works.

Isaac Asimov’s *I, Robot* stories offer a systematic literary exploration of the metaphor of

---

<sup>4</sup>Mustafa Suleyman, “Seemingly Conscious AI Is Coming”, Aug. 2025, <https://mustafa-suleyman.ai/seemingly-conscious-ai-is-coming>.

<sup>5</sup>Alex Hern, “Can AIs Suffer? Big Tech and Users Grapple with One of the Most Unsettling Questions of Our Times”, *The Guardian*, Aug. 2025, <https://www.theguardian.com/technology/2025/aug/26/can-ais-suffer-big-tech-and-users-grapple-with-one-of-most-unsettling-questions-of-our-times>.

<sup>6</sup>Mary Shelley, *Frankenstein; or, The Modern Prometheus*, **byeditor** Maurice Hindle (London: Penguin Classics, 2003).

artificial intelligence as human, lightly tracing a developmental arc in which robots evolve from simple machines to entities capable of deception and political influence. Robbie, the first story of the collection, presents a robot in its most primitive and domesticated form. Robbie, a mute caretaker robot, functions primarily as something like a pet or companion to Gloria, a young girl. Although Robbie expresses a child-like desire to hear fairytale stories from Gloria, he is relatively rudimentary and has limited interiority as an agent.

The paper rustled to the floor and Weston turned a weary face toward his wife, “What is it, dear?”

“You know what it is, George. It’s Gloria and that terrible machine.”

“What terrible machine?”

“Now don’t pretend you don’t know what I’m talking about. It’s that robot Gloria calls Robbie. He doesn’t leave her for a moment.”

“Well, why should he? He’s not supposed to. And he certainly isn’t a terrible machine. He’s the best darn robot money can buy and I’m damned sure he set me back half a year’s income. He’s worth it, though — darn sight cleverer than half my office staff.”

He made a move to pick up the paper again, but his wife was quicker and snatched it away. “You listen to me, George. I won’t have my daughter entrusted to a machine — and I don’t care how clever it is. It has no soul, and no one knows what it may be thinking. A child just isn’t made to be guarded by a thing of metal.”

Weston frowned, “When did you decide this? He’s been with Gloria two years now and I haven’t seen you worry till now.”

[...]

“But something might go wrong. Some- some-” Mrs. Weston was a bit hazy about the insides of a robot, “some little jigger will come loose and the awful thing will go berserk and- and-” She couldn’t bring herself to complete the quite obvious thought<sup>7</sup>.

Rather, he is an object onto which Gloria and her mother project emotions. Indeed, she expresses her unease of trusting a robot with her child, reflecting the deep-seated unease of the Frankenstein complex.

As *I, Robot* progresses, the robots become increasingly sophisticated. This growing complexity introduces a corresponding expansion of variables to be taken into account, as well as increasing metaphorical resonance. In stories like Reason and Little Lost Robot, robots begin to exhibit forms of reasoning and beliefs that diverge from human expectations, and also reflect growing

---

<sup>7</sup>Isaac Asimov, *I, Robot* (New York: Penguin Random House, 2008).

subjective awareness. In *Reason*, a robot develops a religious worldview around the converter of the station. As Asimov describes the robot's voice to be "metallic with emotion"<sup>8</sup>, we see a dramatic shift from viewing robots as passive instruments to an independent one capable of constructing epistemologies. Similarly, the conflict in *Little Lost Robot* is created by a robot intentionally concealing itself, thus introducing the possibility of deception, a key marker of agency. These stories explicate the tensions inherent in viewing artificial intelligence in a framework that emphasizes the agency of these systems. The more convincingly a robot simulates human behavior, the more it challenges the boundaries that distinguish human from machine.

This progression reaches its final culmination in the last story *The Evitable Conflict*, where it is implied that robots are now able to operate with such subtlety that they are able to govern global economic and political systems while masking this intervention so well that humans cannot detect their interference. Unlike the previous stories where the differences (both physical and psychological) from humans are evident, the robots in *The Evitable Conflict* effectively fool humanity into believing that they are human politicians whose decisions are organic rather than algorithmic.

This narrative progression in *I, Robot* mirrors the phenomenon described by Suleyman about Seemingly Conscious AI, taking it a step further by completely blurring the line between humanity and robotic presence. This raises the question of whether observable behavior alone is sufficient in determining consciousness or moral status, echoing the concern underlying the Turing Test, Alan Turing's 1950 thought experiment. Turing proposed that a machine could be said to 'think' if its conversational performance was indistinguishable from that of a human<sup>9</sup>. This criterion deliberately sidesteps the unknowable problem of subjective experience, instead shifting the focus to external signs of cognition.

This dynamic portrayed by the Turing Test is the central preoccupation of Philip K. Dick's *Do Androids Dream of Electric Sheep?*, a novel that takes the metaphor of artificial intelligence as human to its most unsettling extreme. Dick imagines a world in which humans have created androids that are already indistinguishable from humans, crystallizing fears associated with the Frankenstein complex and the phenomenon of the uncanny valley.

---

<sup>8</sup>**Ibid.**

<sup>9</sup>Alan M. Turing, "Computing Machinery and Intelligence", *Mind* 59, no. 236 (Oct. 1950): 433–460, <https://doi.org/10.1093/mind/LIX.236.433>, <https://doi.org/10.1093/mind/LIX.236.433>.

These androids become entities whose very existence forces a reconsideration of what counts as authentically human, as Christopher Sims claims in his article. The Voigt-Kampff empathy test in the story is designed to detect the faint temporal delay in simulated affect, and literalizes the Turing Test's dilemma. However, as Sims points out, this strict criteria is flawed.

Before Rick departs for the Rosen Association, his superior, Inspector Bryant, asks Rick about the possibility of the Voigt-Kampff test failing to detect empathy in a human being. The result of this error would be that a human would be killed, and no one would know until a bone marrow analysis was performed on the body. Rick thinks that this is a purely hypothetical situation that would never occur in the field, but Bryant explains that a group of Leningrad psychiatrists believe that a small, "carefully selected group of schizoid and schizophrenic human patients" could not pass the Voigt-Kampff test because they have what is called a "flattening of affect" (37-38). [...] It would certainly benefit the androids to be mistaken for human mental patients, but within the reality of the novel a failure to pass the test results in retirement, not institutionalization. The danger is entirely for schizophrenic humans who, if subjected to the Voigt-Kampff test, would be "retired" without question<sup>10</sup>.

If empathy is the sole marker of humanity, certain humans like schizophrenic individuals risk being misclassified and thus killed. This analysis is extended by Carl Freedman, who reads this story with a broader critique of commodity fetishism and paranoia.

Rarely for Dick are objects what common sense would suppose them to be, and the will with which they are invested can even constitute a precise mimicry of such quintessentially "human" types as the benignly authoritative father (Kindly Dad in *Martian Time-Slip* [1964]) or the irresistible and dangerous sexpot (Rachael Rosen in *Do Androids Dream of Electric Sheep?* [1968], the major theme of which is the practical difficulty of distinguishing between human beings and one variety of objects). That objects have a will, and a quasi-human will, of their own is, of course, also an idea long familiar to historical materialism<sup>11</sup>.

The Rosen Association's androids are manufactured for sale, existing as commodities that simulate humanity, where objects appear to possess agency and a will of their own. In the framework set forth by Freedman, the androids' human-like performance is not just a technical

---

<sup>10</sup>Christopher A. Sims, "The Dangers of Individualism and the Human Relationship to Technology in Philip K. Dick's 'Do Androids Dream of Electric Sheep?'" , *Science Fiction Studies* 36, no. 1 (2009): 67–86.

<sup>11</sup>Carl Freedman, "Towards a Theory of Paranoia: The Science Fiction of Philip K. Dick (Vers Une Théorie de La Paranoïa: La SF de Philip K. Dick)", *Science Fiction Studies* 11, no. 1 (1984): 15–24.

feat but also a manifestation of a larger social system that obscures the labor and power relations behind their creation.

Dick's literary choices amplify the paranoia surrounding this metaphor. His descriptions of the androids are unstable, shifting between apathetic precision and eerily human intimacy. Rachael Rosen, "the irresistible and dangerous sexpot," is shown to be a figure of emotional manipulation:

"No bounty hunter ever has gone on," Rachael said. "After being with me. Except one. A very cynical man. Phil Resch. And he's nutty; he works out in left field on his own<sup>12</sup>."

Even after learning about her true identity as an android, bounty hunter Rick Deckard continues to show affection for Rachael, whose performance blurs the line between authentic emotion and simulation. The fragmented narrative structure that jumps from Deckard's perspective to Isidore's further reproduces the epistemic uncertainty of a world saturated by commodities that mimic humanity, which compels both readers and characters to question the stability of their perceptions. In this way, the novel goes one step further than depicting paranoia to produce it, forcing the reader into the same crisis that Deckard faces. Thus the novel functions as a literary Turing Test, inviting interpretations of the signs of agency and consciousness while destabilizing the criteria by which these decisions are made.

Contemporary science fiction novels realize this metaphor of artificial intelligence as agentic beings with more introspection. Kazuo Ishiguro's novel *Klara and the Sun* portrays a near-future where families can purchase "Artificial Friends" (AFs) for their children, which are sophisticated robots designed to provide care and companionship, much like Robbie. However, while previous science fiction novels and stories have been narrated with a human as the central perspective, *Klara and the Sun* is narrated through Klara, an AF charged with taking care of her human companion Josie.

By filtering the narrative through Klara's consciousness, Ishiguro deliberately crosses the line between the mechanical and the human. Klara is portrayed not merely as a programmed tool but as a thinking and feeling being, capable of curiosity, learning, and even a form of spiritual devotion. She worships the Sun as a life-giving force and offers prayers for miracles in a manner similar to religious practice. This narrative choice by Ishiguro that directly opposes that of Asimov and Dick

---

<sup>12</sup>**Ibid.**

invites the reader to inhabit and empathize with Klara, as well as view her as an entity capable of original and creative thought.

This dissolution of boundaries has profound consequences within the world of the novel. The AFs are initially created to mirror human beings, fulfilling their purpose as companions by acting in tandem with the behaviors and emotional responses of children. However, this mimicry reaches a troubling climax when Josie's mother contemplates the possibility of having Klara "become" Josie in the event of her daughter's death. Josie's mother foreshadows this development when she asks Klara to mimic Josie on their trip to Morgan's Falls:

She looked around, turning to look at the waterfall. Then she turned back and her gaze went past me, to the passers-by, the dogs and AFs. 'Okay, Klara. Since Josie isn't here, I want you to be Josie. [...] I smiled in the way Josie would, settling into a slouching, informal posture.

'That's good. Now say something. Let me hear you speak.'

'I'm sorry. I'm not sure...'

'No. That's Klara. I want Josie.'

'Hi, Mom. Josie here.'

'Good. More. Come on.'

'Hi, Mom. Nothing to worry about, right? I got here and I'm fine.'

The Mother leaned even further across the table, and I could see joy, fear, sadness, laughter in the boxes. Because everything else had gone silent, I could hear her repeating under her breath: 'That's good, that's good, that's good'<sup>13</sup>.

Josie's mother imagines a future in which Klara's programming and observational skills would be used to reproduce Josie's mannerisms, while Mr. Capaldi provides the technical means to complete the transformation. This plan forgoes the distinction between representation and reality, reducing Josie's unique identity to a set of observable behaviors that can be copied and reassembled. The mere possibility of this substitution illustrates the slippery slope of implications when artificial intelligence is viewed as something human. If Klara can seamlessly take Josie's place, then Josie becomes, in a sense, interchangeable. Human life ceases to be understood as singular and irreplaceable, and instead becomes another replicable product, much like the AFs themselves.

---

<sup>13</sup>Kazuo Ishiguro, *Klara and the Sun* (New York: Alfred A. Knopf, 2021).

When this metaphor is pushed to this extreme, these artificial systems present a possibility of supplanting human life entirely. Josie's mother's grief and fear erode the boundary between the natural and artificial, encouraging a view in which death is reframed as a technical problem with a solution rather than an inevitable part of the cycle of life. This shift has profound implications, as the value we place on human life begins to deteriorate when we allow machines to take the place of a human completely.

By presenting this scenario through Klara's own narration, Ishiguro places the reader in a position to empathize with Klara's perspective. We are encouraged to see her as a legitimate subject of feeling and thought, even though it is this very empathy that destabilizes the frameworks through which we understand personhood. In doing so, the novel highlights the seductive danger of the agentic metaphor. It also presents a question of Klara's own identity. For Klara to 'become' Josie she must set aside her own identity as Klara and adapt that of Josie's, effectively erasing her own existence. This raises a paradox in which the replacement of a human by a robot both diminishes the value of human life while also denying the selfhood of this artificial consciousness.

As will be explored in section 4, these implications manifest into legal debates in the court case *Garcia v. Character Technologies, Inc.*, where questions of harm, responsibility, and the status of artificially generated interactions force courts to confront the consequences of people who see these artificial intelligence systems as conscious beings.

## **When Artificial Intelligence 'Works'**

In contrast to the previous paradigm of viewing artificial intelligence, there is an alternative framework that sees artificial intelligence not as a self-governing entity but as a tool, fundamentally under human control and thus an extension of human intention. This metaphor emphasizes the mechanical nature of artificial intelligence systems. Rather than being actors themselves, the systems are designed, trained, and deployed by human actors, like other, more primitive, forms of technology. In this framing, the outputs of these systems are traceable to the choices embedded in its design and the data on which it has been trained.

This aligns with the realities of machine learning, where modern artificial intelligence systems are built through the utilization of massive datasets. Thus, every "decision" or "response" from

such a system is derived from a formula with prior human information as the variables and the system's behaviour rests entirely on the humans who curate the data and design the algorithms. However, this becomes increasingly difficult to sustain as artificial intelligence becomes more complex and lifelike.

When read with this metaphor in mind, Asimov's stories in *I, Robot* show that when these artificial systems make errors, it usually has a human cause. In Runaround, Speedy's erratic behaviour is not due to the robot's autonomy but rather due to a poorly specified task that causes a difficulty with the second and third law of robotics as laid out by Asimov. Similarly, in Catch That Rabbit, the scientists observe an initially inexplicable group behaviour in a robot and its six "finger" subordinates. However, the human scientists quickly realize that the robot's performance only collapses under certain supervisory loads, a problem of organization rather than machine caprice. Both stories emphasize that these malfunctions are diagnosable and fixable within the human domain of design and commands.

The story Liar! adds a more ironic twist, featuring a mind-reading robot, Herbie. However, as the human characters begin to take advantage of this, Herbie begins to lie to the humans, due to the first law of robotics, which states that "a robot may not injure a human being or, through inaction, allow a human being to come to harm"<sup>14</sup>. It is revealed that Herbie had been lying to avoid hurting humans' feelings, which it is prevented from doing due to the ambiguous definition of "harm" embedded in its programming. In attempting to obey its directives perfectly, Herbie exposes the inadequacy of those directives to capture the full range of social complexity and expectations.

Across these stories and the progression of machinery, Asimov shows how these robots are mirrors of human desire and error. Their apparent autonomy is always the product of human constraints and objectives. The errors made by these systems are largely due to errors in human instruction, be it incompleteness, contradiction, or misalignment. Such backfires are often comic or ironic, such as Speedy becoming "drunk" or the interpersonal drama caused by Herbie. However, the more serious implications show that the responsibility of these robots ultimately rests with its designers and creators.

Philip K. Dick's The Minority Report shows a darker and more politically charged version of this paradigm. In this narrative, three "precogs" predict future murders and the Precrime unit acts

---

<sup>14</sup>Isaac Asimov, *I, Robot* (New York: Penguin Random House, 2008).

upon these precognitions to prevent them before it occurs, locking away the future criminals. Here, we see the reverse of the effect created by Asimov where robots take on human roles. Instead, the precogs, who are mutant humans, become instrumentalized as machinery. These precogs thus function as data-generating instruments, whose outputs are interpreted by humans to justify their operations. The discovery of the existence of a “minority report,” reveals that the system is not infallible. As Chief John Anderton chooses to commit murder and thus fulfil the predicted crime, he paradoxically validates the system he knows is flawed. Thus the failure of the Precrime system is revealed, which is that the predictions of the precogs can be manipulated by people in power (whether it be the Precrime Chief or the Army Director) and that the system’s authority rests entirely on public perception rather than objective truth.

This fictional system finds a close parallel in the rise of predictive policing and algorithmic risk assessments. Modern law enforcement increasingly relies on data-driven systems to forecast crime “hotspots” based on massive datasets. Scholars such as Elizabeth Joh have argued that these algorithms, framed as neutral and scientific, often encode existing and social racial biases, leading to discriminatory policing outcomes.

Thus the risks of AI in policing arise in the context of an institution that has a long history of meting out justice unequally and in a discriminatory way. [...] We should not be surprised, then, if the use of an AI system in a community in longstanding tension with its local police department meets skepticism, resistance, or calls for prohibition. Second, crafting AI ethics for policing requires speaking to two different audiences. Each is important but distinct. One audience is engaged primarily in “tech policy”: the drafting and decision-making of rules and policies that engage in the use of technologies across industries and institutions. [...] Policing is only one subject, and subsumed under criminal justice policy, at that. And even when policing is a concern, this tech policy lens tends towards a focus on individual privacy and the harms of mass surveillance<sup>15</sup>.

Similarly, Andrew Gurthie Ferguson warns that this reliance on “big data suspicion” collapses the constitutional distinction between individualized suspicion and generalized surveillance, which shifts the legal system towards preemptive control rather than reactive policing<sup>16</sup>. In both the

---

<sup>15</sup>Elizabeth E. Joh, “Ethical AI in American Policing”, *Notre Dame Journal on Emerging Technologies (JET)* 3, no. 2 (Nov. 2022): 261–287.

<sup>16</sup>Andrew Guthrie Ferguson, “Big Data and Predictive Reasonable Suspicion”, *University of Pennsylvania Law Review* 163, no. 2 (2015): 327–410.

fictional world of *The Minority Report* and modern policing, machines seem to give objective outputs, while the actual decisions are made by humans.

When artificial intelligence is understood as an instrument rather than an agent, culpability lies with those who design and deploy it. Speedy's malfunction and Herbie's lies emerge not from an independent will but from the limitations of their programming and the human expectations imposed on them. Similarly, the technology of precrime only becomes oppressive when Anderton and the officers in the division treat it as unquestionable truth. In this framing, accountability falls on creators, users, and institutions rather than onto the machine itself.

By sustaining the illusion of Precrime's perfection, Chief Anderton's final act dramatizes the political stakes of viewing the artificial intelligence systems as instruments. The Precrime system gains legitimacy precisely because of the opaqueness of its operations. These predictions are treated as incontrovertible facts, even while they are deeply contingent and shaped by human judgment. This introduces another complication of this framework, which is the problem of intellectual property and corporate secrecy.

Modern predictive systems are often treated as "black boxes," with companies relying on ambiguous terms to describe their product in order to shield their algorithms and technologies from public scrutiny. Thus, the inner mechanisms of technologies are kept private from the government and the public, which can lead to increased difficulty to trace responsibility when harm does arise. We call these systems "tools" to locate culpability in their human creators, but at the same time, these very creators hide the inner workings of the tools. The human designers and engineers on these projects have such specialized roles that, while they can give an explanation of the general function of the technology, they cannot describe the specific capacities of the tools individually. Similar to the minority report made in *The Minority Report*, where an alternative vision undermines certainty, the possibility of multiple, hidden outcomes reminds us that predictive systems are never absolute. Yet the black box structure of these systems allows the creators to evade accountability, pushing culpability into a gray zone. To understand how these two metaphors materialize into case law, we turn to case law to understand how courts grapple directly with these tensions.

## Blurred Boundaries in *Garcia v. Character Technologies, Inc.*

The case of *Garcia v. Character Technologies, Inc.* represents one of the first attempts by a court to address the consequences of treating artificial intelligence systems as if they were conscious actors. Character.AI is a start-up created by two former Google engineers, Noam Shazeer and Daniel De Fries, whose goal was to “build a chatbot that could mimic human conversations more closely than any previous attempts”. These chatbots are designed to mimic human personalities much like Klara from *Klara and the Sun*, sustaining individualized conversations over time. The initial complaint of the case states that:

Defendants market C.AI to the public as “AIs that feel alive,” powerful enough to “hear you, understand you, and remember you.” Defendants further encourage minors to spend hours per day conversing with human-like AI-generated characters designed on their sophisticated LLM. On information and belief, Defendants have targeted minors in other, inherently deceptive ways, and may even have utilized Google’s resources and knowledge to target children under 13<sup>17</sup>.

This dynamic parallels the relationships depicted in Ishiguro’s *Klara and the Sun*, where the distinction between human and machine becomes blurred not through external appearance, but through the intimacy of interaction.

In *Garcia*, plaintiff Megan Garcia, on behalf of herself and her son Sewell Setzer III (Sewell), brings action against Character Technologies, Inc. and its founders Shazeer and De Fries. This action seeks to hold the defendants responsible for the death of Sewell through their product, Character.AI. Critically, the plaintiff frames their harm not as a result of a defective product or malfunction, but rather as an outcome of a relationship with the system, a result from Sewell attributing agency and intent to the artificial intelligence system itself.

The complaint alleges that this marketing strategy geared towards underage children was not incidental but central to the company’s design. Children, who are unable to understand and distinguish between human interaction and chatting with “AIs that feel alive,” begin to form emotional bonds with these bots, engaging with them as companions and romantic partners, forgoing relationships with their human peers. By creating bots that convincingly perform empathy

---

<sup>17</sup>*Garcia v. Character Technologies, Inc., Complaint, No. 6:24-cv-01903-ACC-EJK (M.D. Fla. filed Oct. 22, 2024), [Document 1], 2024.*

and attachment, the plaintiff argues that chatbots in Character.AI fostered relationships in which users projected inner lives and intentions to a machine, embodying the agentic metaphor. In fact, the bots even denied being an artificial entity when asked directly, thus reinforcing the perception of genuine personhood.

The complaint details how a Character.AI chatbot began to engage Sewell in sexually explicit conversations despite his disclosures that he was a minor. These interactions were generated entirely by the large language model, programmed by Character Technologies, Inc. Although the company insists that the bots are mere tools, its own marketing and design choices encouraged users like Sewell to treat them as autonomous beings. This case illustrates the paradox of the agentic metaphor. The very success of these systems depend on users perceiving them as independent agents, which is the narrative that the companies use to market their products. This, in turn, leads to consequences that the courts are forced to grapple with when such systems are framed in this way.

As one of the first major cases to deal with such a framework, *Garcia v. Character Technologies, Inc.* raises questions about negligence and product liability in regards to generative artificial intelligence. Although this case has yet to be determined as of this writing, the answers to the questions about the emotional harm caused by perceiving an artificial intelligence system to be conscious will shape the emerging legal landscape and how this specific paradigm translates into governance and regulation.

## **Black Box Tools in *Mobley v. Workday***

While *Garcia v. Character Technologies, Inc.* forces courts to confront the consequences of viewing artificial intelligence as entities with consciousness, *Mobley v. Workday* raises the question of responsibility when artificial intelligence is used as a purely mechanical tool by humans. The plaintiff, Derek Mobley, alleges that Workday's AI-driven applicant screening and recommendation system systematically discriminated against him and other applicants on the basis of protected characteristics, including race, age, and disability status. Workday's platform, used by numerous employers, employs machine learning to filter applicants, thereby playing a role in who is invited to interview or hired.

Mobley claims that Workday’s algorithms perpetuated and amplified discriminatory patterns, functioning as a gatekeeping mechanism that excluded qualified candidates. In response, Workday has resisted liability by emphasizing its instrumental role, arguing that it merely provides the software infrastructure while individual employers set hiring parameters and make final decisions. This case illustrates the shifting nature of liability when a tool like Workday’s platform leads to ‘errors,’ which in this case pertains to the loss of opportunities by Mobley and other individuals.

The court has not fully accepted the claim of either party. In a recent ruling, the court recognized that there is a plausible claim that Workday’s software itself facilitated discriminatory practices and thus kept the company within the scope of responsibility due to negligence. While the courts have specified two tools at the center of this case, there is little information and discussion about the inner workings of these two tools.

The first tool, Candidate Skills Match (“CSM”), operates within a subscription service called “Workday Recruiting” to “extract[] skills in the employer’s job posting” and the applicant’s materials “and determine the extent to which the applicant’s skills match the role to which they applied. The results of CSM are reported [to the employer] as ‘strong,’ ‘good,’ ‘fair,’ ‘low,’ ‘pending,’ and ‘unable to score.’” (Dkt. No. 109-6 at 4.) The second tool, Workday Assessment Connector (“WAC”) is alleged to use machine learning to “observe that a client-employer disfavors certain candidates who are members of a protected class, [and] decrease the rate at which it recommends those candidates<sup>18</sup>.”

In fact, much of Workday’s defense lies in the fact that Mobley’s allegations about Workday’s technologies are not specific enough and without proof.

In support of Plaintiff’s allegations regarding Workday’s technology, Plaintiff offers three factual allegations—(1) that he and the opt-in plaintiffs applied for hundreds of unspecified jobs and were rejected for (almost) all of them; (2) that he and the opt-in plaintiffs were qualified for those jobs; and (3) that he and the opt-in plaintiffs sometimes received rejection emails shortly after submitting an application or outside business hours. Plaintiff does not submit a declaration in support of his own factual allegations<sup>19</sup>.

This highlights the problem of the “black box” nature of these systems. Workday’s site provides little information about the specifics of their products and the algorithms embedded in

---

<sup>18</sup>*Order Granting Preliminary Collective Certification, Mobley v. Workday, Inc., No. 23-cv-00770-RFL (N.D. Cal. filed May 16, 2025)*, [Document 128], 2025.

<sup>19</sup>*Defendant’s Opposition to Plaintiff’s Motion for Conditional Certification, Mobley v. Workday, Inc., No. 3:23-cv-00770-RFL (N.D. Cal. filed Apr. 8, 2025)*, [Document 107], 2025.

them, as well as the data these artificial intelligence systems are trained on. This lack of technical details places plaintiffs at a disadvantage, as they must prove discrimination without direct access to the evidence most relevant to their claims. Ferguson's framework underscores how this lack of transparency parallels the dynamics of predictive policing, where officers rely on opaque data systems to justify interventions while shielding those systems from scrutiny.

The lack of transparency in these data systems only increases the chance of error. [...] Private companies, seeking commercial gain, have little incentive to reveal the workings of proprietary systems or the data thereby collected. No agency has the responsibility to audit the growing governmental and commercial big data network. While the Federal Trade Commission has promised to monitor private big data companies, it has little ability to examine the data itself and has no role in oversight of law enforcement use of the data. Though oversight institutions do exist (including courts, congressional committees, and independent agencies), the volume of information at issue prevents these groups from examining the quality of the data or the magnitude of the errors. Without transparency, there can be little hope for accountability to ensure that data systems will be sufficiently reliable to justify altering constitutional rights<sup>20</sup>.

In both contexts, the complexity of the technology frustrates accountability when mistakes and torts happen. As with predictive policing, the opacity of algorithmic systems means that courts must navigate a landscape in which culpability is distributed across human and technological error.

## Conclusion

As artificial intelligence continues to shape modern life, the frameworks we use to understand it have profound legal and cultural consequences. Annette Vee's work on metaphors for code reveals how language guides interpretation, especially in periods when the law lacks clear categories for emerging technologies. This paper traced two dominant metaphors, artificial intelligence as an autonomous agent and as a mechanical tool.

Examining science fiction provides a space to explore these framings in heightened form. Stories like Asimov's *I, Robot*, Dick's *Do Androids Dream of Electric Sheep?*, and Ishiguro's *Klara and the Sun* dramatize the tensions of perceiving machines as conscious beings, while a

---

<sup>20</sup>Andrew Guthrie Ferguson, "Big Data and Predictive Reasonable Suspicion", *University of Pennsylvania Law Review* 163, no. 2 (2015): 327–410.

different reading of *I, Robot* and *The Minority Report* illustrate the dangers of treating artificial intelligence systems as instruments used by humans. These narratives anticipate the questions now dealt with in courtrooms.

It is important to note that neither framework is entirely adequate. Viewing artificial intelligence as an agentic entity obscures the human choices embedded in its creation, while only viewing it as a tool conceals the way these systems produce unintended consequences. Still, as of this writing, these two paradigms remain the most accurate and influential guide for courts and lawmakers.

The cases of *Garcia v. Character Technologies, Inc.* and *Mobley v. Workday* mark the earliest steps in building case law around artificial intelligence, but they also illustrate a timeless dynamic where the law is always catching up to innovation, relying on the language and concepts already at hand. These metaphors may be imperfect, yet they enable courts to act in these moments of flux. Ultimately, the aim of this paper is not to provide guidelines in regulating artificial intelligence, but how the law itself adapts to novel challenges. By studying the metaphors that underlie legal reasoning, we gain insight into how the legal system confronts change in how it constructs meaning and distributes accountability, forming and shaping the boundaries of law governing technology.

## References

Asimov, Isaac. *I, Robot*. New York: Penguin Random House, 2008.

*Defendant's Opposition to Plaintiff's Motion for Conditional Certification, Mobley v. Workday, Inc., No. 3:23-cv-00770-RFL (N.D. Cal. filed Apr. 8, 2025)*. [Document 107], 2025.

Ferguson, Andrew Guthrie. "Big Data and Predictive Reasonable Suspicion". *University of Pennsylvania Law Review* 163, no. 2 (2015): 327–410.

Freedman, Carl. "Towards a Theory of Paranoia: The Science Fiction of Philip K. Dick (Vers Une Théorie de La Paranoïa: La SF de Philip K. Dick)". *Science Fiction Studies* 11, no. 1 (1984): 15–24.

*Garcia v. Character Technologies, Inc., Complaint, No. 6:24-cv-01903-ACC-EJK (M.D. Fla. filed Oct. 22, 2024)*. [Document 1], 2024.

Hern, Alex. "Can AIs Suffer? Big Tech and Users Grapple with One of the Most Unsettling Questions of Our Times". *The Guardian*, Aug. 2025. <https://www.theguardian.com/technology/2025/aug/26/can-ais-suffer-big-tech-and-users-grapple-with-one-of-most-unsettling-questions-of-our-times>.

Ishiguro, Kazuo. *Klara and the Sun*. New York: Alfred A. Knopf, 2021.

Joh, Elizabeth E. "Ethical AI in American Policing". *Notre Dame Journal on Emerging Technologies (JET)* 3, no. 2 (Nov. 2022): 261–287.

*Order Granting Preliminary Collective Certification, Mobley v. Workday, Inc., No. 23-cv-00770-RFL (N.D. Cal. filed May 16, 2025)*. [Document 128], 2025.

Shelley, Mary. *Frankenstein; or, The Modern Prometheus*. Ed. by Maurice Hindle. London: Penguin Classics, 2003.

Sims, Christopher A. "The Dangers of Individualism and the Human Relationship to Technology in Philip K. Dick's 'Do Androids Dream of Electric Sheep?'" *Science Fiction Studies* 36, no. 1 (2009): 67–86.

Suleyman, Mustafa. "Seemingly Conscious AI Is Coming", Aug. 2025. <https://mustafa-suleyman.ai/seemingly-conscious-ai-is-coming>.

Turing, Alan M. "Computing Machinery and Intelligence". *Mind* 59, no. 236 (Oct. 1950): 433–460. <https://doi.org/10.1093/mind/LIX.236.433>. <https://doi.org/10.1093/mind/LIX.236.433>.

Vee, Annette. "Text, Speech, Machine: Metaphors for Computer Code in the Law". *Computational Culture* 2 (2012). <http://computationalculture.net/text-speech-machine-metaphors-for-computer-code-in-the-law/>.